Storypoint Problem - appceleratorstudio

September 5, 2024

1 Storypoint Prediction: Regression Approach

1.1 Preparation

```python
import os
import json

import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import seaborn as sns
from scipy.sparse import csr_matrix, hstack, vstack

from sklearn.pipeline import Pipeline
from sklearn.preprocessing import RobustScaler
from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error,
 ↪f1_score, precision_score, recall_score, accuracy_score
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.model_selection import cross_val_score, learning_curve,
 ↪validation_curve
from trainer import GridSearchCVTrainer

project_name = 'appceleratorstudio'
```

1.1.1 Plot learning curve

```python
def plot_learning_curve(estimator, title, X, y, ylim=None, cv=None,
                        n_jobs=-1, train_sizes=np.linspace(.1, 1.0, 10)):
    plt.figure()          # Create new figure
    plt.title(title)      # Set title of the figure

    # Set y-axis limits: ylim=(min, max)
    if ylim is not None:
        plt.ylim(*ylim)

    plt.xlabel("Training examples")   # Set x-axis label
    plt.ylabel("Score")               # Set y-axis label
```

```python
    # Generate learning curve data
    train_sizes, train_scores, test_scores = learning_curve(
        estimator, X, y, cv=cv, n_jobs=n_jobs, train_sizes=train_sizes,␣
↪scoring='neg_mean_squared_error')

    train_scores_mean = np.mean(train_scores, axis=1)  # Calculate mean of␣
↪training scores
    train_scores_std = np.std(train_scores, axis=1)    # Calculate standard␣
↪deviation of training scores
    test_scores_mean = np.mean(test_scores, axis=1)    # Calculate mean of test␣
↪scores
    test_scores_std = np.std(test_scores, axis=1)      # Calculate standard␣
↪deviation of test scores

    plt.grid()  # Display grid

    # Fill the area between the mean training score and the mean +/- std␣
↪training score
    plt.fill_between(train_sizes, train_scores_mean - train_scores_std,
                     train_scores_mean + train_scores_std, alpha=0.1,
                     color="r")

    # Fill the area between the mean test score and the mean +/- std test score
    plt.fill_between(train_sizes, test_scores_mean - test_scores_std,
                     test_scores_mean + test_scores_std, alpha=0.1, color="g")

    # Plot mean training score as points
    plt.plot(train_sizes, train_scores_mean, 'o-', color="r",
             label="Training score")
    # Plot mean test score as points
    plt.plot(train_sizes, test_scores_mean, 'o-', color="g",
             label="Validation score")

    plt.legend(loc="best")  # Display legend
    return plt
```

1.1.2 Plot validation curve

```python
[ ]: def plot_validation_curve(estimator, title, X, y, param_name, param_range,
                               y_lim=None, cv=10, n_jobs=-1):
         train_scores, val_scores = validation_curve(estimator=estimator, X=X, y=y,
                                                     param_name=param_name,␣
↪param_range=param_range,
                                                     cv=cv, n_jobs=n_jobs,
                                                     ␣
↪scoring='neg_mean_squared_error')
```

```python
    # Calculate mean and standard deviation of training and validation scores
    train_mean = np.mean(train_scores, axis=1)
    tran_std = np.std(train_scores, axis=1)
    val_mean = np.mean(val_scores, axis=1)
    val_std = np.std(val_scores, axis=1)
    print(val_mean)

    # Plot train scores
    plt.plot(param_range, train_mean, color='r', marker='o', markersize=5,
 label='Training score')
    plt.fill_between(param_range, train_mean + tran_std, train_mean - tran_std,
 alpha=0.15, color='r')

    # Plot validation scores
    plt.plot(param_range, val_mean, color='g', linestyle='--', marker='s',
 markersize=5, label='Validation score')
    plt.fill_between(param_range, val_mean + val_std, val_mean - val_std,
 alpha=0.15, color='g')

    plt.title(title)          # Set title of the plot
    plt.grid()                # Display grid
    plt.xscale('log')         # Set x-axis scale to log
    plt.legend(loc='best')    # Display legend
    plt.xlabel('Parameter')   # Set x-axis label
    plt.ylabel('Score')       # Set y-axis label

    # Set y-axis limits
    if y_lim != None:
        plt.ylim(y_lim)

    return plt
```

1.1.3 Evaluate model

```python
[ ]: def evaluate_model(model, model_name, X_test, y_test, y_logscale=False,
 save_directory=None):
    y_pred = model.predict(X_test)
    if(y_logscale):
        y_pred = np.exp(y_pred)

    lines = [model_name + '\'s evaluation results:']

    mse = mean_squared_error(y_test, y_pred)
    rmse = np.sqrt(mse)
    mae = mean_absolute_error(y_test, y_pred)
    r2 = r2_score(y_test, y_pred)
```

```python
    lines.append(f' - Mean squared error:      {mse:.2f}')
    lines.append(f' - Root mean squared error: {rmse:.2f}')
    lines.append(f' - Mean absolute error:     {mae:.2f}')
    lines.append(f' - R2 error:                {r2:.2f}')

    y_pred = np.round(y_pred).astype(int)
    f1 = f1_score(y_test, y_pred, average='weighted')
    precision = precision_score(y_test, y_pred, average='weighted',␣
 ↪zero_division=0)
    recall = recall_score(y_test, y_pred, average='weighted', zero_division=0)
    accuracy = accuracy_score(y_test, y_pred)

    lines.append(f' - F1 score:                {f1:.2f}')
    lines.append(f' - Precision:               {precision:.2f}')
    lines.append(f' - Recall:                  {recall:.2f}')
    lines.append(f' - Accuracy:                {accuracy:.2f}')
    lines.append('---------------------------------------------')
    lines.append('')

    # Save to file

    if(save_directory != None):
        filename = save_directory + project_name + '.txt'
        directory = os.path.dirname(filename)
        if not os.path.exists(directory):
            os.makedirs(directory)
        with open(filename, 'a') as f:
            for line in lines:
                print(line)
                f.write(line + '\n')
    else:
        for line in lines:
            print(line)
```

1.1.4 Set random seed

```python
import numpy as np
import random
import os

# Set random seed for numpy
np.random.seed(42)

# Set random seed for random
random.seed(42)

# Set random seed for os
```

```
os.environ['PYTHONHASHSEED'] = '42'
```

1.2 Dataset set-up

1.2.1 Bag of Words preprocessing

This is a Bag of Words preprocess approach. I will use 2 CountVectorizer from sklearn to change title and description to two 2 vectors and then concatenate them together. In the rest of this notebook, I will use cross-validation instead hold-out. Therefore, I will join the validation set with training set.

NOTE: I don't add the length of description (and the title) because it has a big deviation which may cause noise.

```python
[ ]:   # Import and remove NaN value


data_train = pd.concat([pd.read_csv('data/' + project_name + '/' + project_name␣
 ↪+ '_train.csv'),
                        pd.read_csv('data/' + project_name + '/' + project_name␣
 ↪+ '_valid.csv')])
data_test = pd.read_csv('data/' + project_name + '/' + project_name + '_test.
 ↪csv')

data_train['description'].replace(np.nan, '', inplace=True)
data_test['description'].replace(np.nan, '', inplace=True)

# Vectorize title
title_vectorizer = CountVectorizer(ngram_range=(1, 2), min_df=2)
title_vectorizer.fit(pd.concat([data_train['title'], data_test['title']]))

# Vectorize description
description_vectorizer = CountVectorizer(ngram_range=(1, 2), min_df=2)
description_vectorizer.fit(pd.concat([data_train['description'],␣
 ↪data_test['description']]))


X_train = hstack([title_vectorizer.transform(data_train['title']).astype(float),
                  description_vectorizer.transform(data_train['description']).
 ↪astype(float),
                  data_train['title'].apply(lambda x : len(x)).to_numpy().
 ↪reshape(-1, 1),
                  data_train['description'].apply(lambda x : len(x)).to_numpy().
 ↪reshape(-1, 1)
                 ])

y_train = data_train['storypoint'].to_numpy().astype(float)
```

```python
X_test = hstack([title_vectorizer.transform(data_test['title']).astype(float),
                 description_vectorizer.transform(data_test['description']).
  ↪astype(float),
                 data_test['title'].apply(lambda x : len(x)).to_numpy().
  ↪reshape(-1, 1),
                 data_test['description'].apply(lambda x : len(x)).to_numpy().
  ↪reshape(-1, 1)
                ])

y_test = data_test['storypoint'].to_numpy().astype(float)
```

```python
print('Check training dataset\'shape:', X_train.shape, y_train.shape)
print('Check testing dataset\'shape:', X_test.shape, y_test.shape)
```

```
Check training dataset'shape: (2589, 21705) (2589,)
Check testing dataset'shape: (287, 21705) (287,)
```

I will use log-scale the label to get a normal distribution of it.

```python
y_train_log = np.log(y_train)
```

1.3 Model training

1.3.1 Linear Regressor

```python
from sklearn.linear_model import ElasticNet
```

Define params-grid:

```python
dict_param = {
    'alpha': [.0001, .001, .01, .1, 1, 10, 100, 1000, 10000],
    'l1_ratio': [.0, .2, .4, .6, .8, 1],
    'max_iter': [2000],
    'random_state': [42]
}
```

```python
gridsearch = GridSearchCVTrainer(name='Elastic Net', model=ElasticNet(),
  ↪param_grid=dict_param, cv=5, n_jobs=5)
gridsearch.load_if_exists()
gridsearch.fit(X_train, y_train_log)

elastic_model = gridsearch.best_estimator_
elastic_model.fit(X_train, y_train_log)
```

```
0it [00:00, ?it/s]
```

```
ElasticNet(alpha=100, l1_ratio=0.4, max_iter=2000, random_state=42)
```

```
evaluate_model(elastic_model, 'Elastic Net model', X_test, y_test,␣
 ↪y_logscale=True, save_directory='results/BoW/')
```

```
Elastic Net model's evaluation results:
 - Mean squared error:      3.69
 - Root mean squared error: 1.92
 - Mean absolute error:     1.33
 - R2 error:                -0.00
 - F1 score:                0.34
 - Precision:               0.26
 - Recall:                  0.51
 - Accuracy:                0.51
-------------------------------------------------
```

```
elastic_model.get_params()
```

```
{'alpha': 100,
 'copy_X': True,
 'fit_intercept': True,
 'l1_ratio': 0.4,
 'max_iter': 2000,
 'positive': False,
 'precompute': False,
 'random_state': 42,
 'selection': 'cyclic',
 'tol': 0.0001,
 'warm_start': False}
```

1.3.2 Support Vector Regressor

```python
from sklearn.svm import SVR
```

```python
dict_param = {
    'C': [.0001, .001, .01, .1, 1, 10, 100, 1000, 10000],
    'gamma': np.logspace(-9, 3, 13),
    'kernel': ['rbf']
}
```

```python
grid_search = GridSearchCVTrainer(name="Support Vector Regressor",model=SVR(),␣
 ↪param_grid=dict_param, cv=5, n_jobs=5)
grid_search.load_if_exists()
grid_search.fit(X_train, y_train_log)

svr_model = grid_search.best_estimator_
svr_model.fit(X_train, y_train_log)
```

There is no checkpoint file for this model.

```
100%|        | 117/117 [05:35<00:00,  2.87s/it]
```

[ ]: SVR(C=1000, gamma=1e-06)

[ ]:
```
evaluate_model(svr_model, 'SVR model', X_test, y_test, y_logscale=True,
 ↪save_directory='results/BoW/')
```

```
SVR model's evaluation results:
 - Mean squared error:      3.71
 - Root mean squared error: 1.93
 - Mean absolute error:     1.50
 - R2 error:               -0.01
 - F1 score:                0.25
 - Precision:               0.31
 - Recall:                  0.23
 - Accuracy:                0.23
--------------------------------------------------
```

[ ]: svr_model.get_params()

[ ]:
```
{'C': 1000,
 'cache_size': 200,
 'coef0': 0.0,
 'degree': 3,
 'epsilon': 0.1,
 'gamma': 1e-06,
 'kernel': 'rbf',
 'max_iter': -1,
 'shrinking': True,
 'tol': 0.001,
 'verbose': False}
```

1.3.3 Random Forest Regressor

[ ]:
```
from sklearn.ensemble import RandomForestRegressor
```

[ ]:
```
dict_param = {
    'max_depth' : [1000, 2000, 5000],
    'min_samples_split': [25, 200, 1000],
    'min_samples_leaf': [1, 2, 3, 4],
    'max_features': [50, 100, 200],
    'n_estimators': [1024],
    'random_state': [42]
}
```

[ ]:
```
grid_search = GridSearchCVTrainer(name="Random Forest Regressor",
                                  model=RandomForestRegressor(),
```

```
                                        param_grid=dict_param, cv = 5, n_jobs=5)
grid_search.load_if_exists()
grid_search.fit(X_train, y_train_log)

rfr_model = grid_search.best_estimator_
rfr_model.fit(X_train, y_train_log)
```

There is no checkpoint file for this model.

100%| | 108/108 [26:20<00:00, 14.64s/it]

```
[ ]: RandomForestRegressor(max_depth=1000, max_features=200, min_samples_split=25,
                           n_estimators=1024, random_state=42)
```

```
[ ]: evaluate_model(rfr_model, 'Random Forest model', X_test, y_test,␣
     ↪y_logscale=True, save_directory='results/BoW/')
```

```
Random Forest model's evaluation results:
 - Mean squared error:      3.37
 - Root mean squared error: 1.83
 - Mean absolute error:     1.35
 - R2 error:                0.09
 - F1 score:                0.31
 - Precision:               0.37
 - Recall:                  0.33
 - Accuracy:                0.33
-------------------------------------------------
```

```
[ ]: rfr_model.get_params()
```

```
[ ]: {'bootstrap': True,
     'ccp_alpha': 0.0,
     'criterion': 'squared_error',
     'max_depth': 1000,
     'max_features': 200,
     'max_leaf_nodes': None,
     'max_samples': None,
     'min_impurity_decrease': 0.0,
     'min_samples_leaf': 1,
     'min_samples_split': 25,
     'min_weight_fraction_leaf': 0.0,
     'monotonic_cst': None,
     'n_estimators': 1024,
     'n_jobs': None,
     'oob_score': False,
     'random_state': 42,
     'verbose': 0,
     'warm_start': False}
```

1.3.4 XGBoost

```python
from xgboost import XGBRegressor
```

```python
dict_param = {
    'eta' : np.linspace(0.01, 0.2, 3),
    'gamma': np.logspace(-2, 2, 5),
    'max_depth': np.asarray([3, 5, 7, 9]).tolist(),
    'min_child_weight': np.logspace(-2, 2, 5),
    'subsample': np.asarray([0.5, .1]),
    'reg_alpha': np.asarray([0.0, 0.05]),
    'n_estimators': np.asarray([10, 20, 50, 100]).tolist(),
    'random_state': [42]
}
```

```python
grid_search = GridSearchCVTrainer(name='XGBoost␣
 ↪Regressor',model=XGBRegressor(), param_grid=dict_param, cv = 5, n_jobs=5)
grid_search.load_if_exists()
grid_search.fit(X_train, y_train_log)

xgb_model = grid_search.best_estimator_
xgb_model.fit(X_train, y_train_log)
```

There is no checkpoint file for this model.

100%| | 4800/4800 [1:03:38<00:00, 1.26it/s]

```
XGBRegressor(base_score=None, booster=None, callbacks=None,
             colsample_bylevel=None, colsample_bynode=None,
             colsample_bytree=None, device=None, early_stopping_rounds=None,
             enable_categorical=False, eta=0.01, eval_metric=None,
             feature_types=None, gamma=0.01, grow_policy=None,
             importance_type=None, interaction_constraints=None,
             learning_rate=None, max_bin=None, max_cat_threshold=None,
             max_cat_to_onehot=None, max_delta_step=None, max_depth=9,
             max_leaves=None, min_child_weight=10.0, missing=nan,
             monotone_constraints=None, multi_strategy=None, n_estimators=100,
             n_jobs=None, num_parallel_tree=None, …)
```

```python
evaluate_model(xgb_model, 'XGBoost Regressor model', X_test, y_test,␣
 ↪y_logscale=True, save_directory='results/BoW/')
```

```
XGBoost Regressor model's evaluation results:
 - Mean squared error:      3.58
 - Root mean squared error: 1.89
 - Mean absolute error:     1.35
 - R2 error:                0.03
 - F1 score:                0.35
 - Precision:               0.37
```

```
    - Recall:                       0.44
    - Accuracy:                     0.44
    ------------------------------------------------
```

[ ]: `xgb_model.get_params()`

[ ]: {'objective': 'reg:squarederror',
    'base_score': None,
    'booster': None,
    'callbacks': None,
    'colsample_bylevel': None,
    'colsample_bynode': None,
    'colsample_bytree': None,
    'device': None,
    'early_stopping_rounds': None,
    'enable_categorical': False,
    'eval_metric': None,
    'feature_types': None,
    'gamma': 0.01,
    'grow_policy': None,
    'importance_type': None,
    'interaction_constraints': None,
    'learning_rate': None,
    'max_bin': None,
    'max_cat_threshold': None,
    'max_cat_to_onehot': None,
    'max_delta_step': None,
    'max_depth': 9,
    'max_leaves': None,
    'min_child_weight': 10.0,
    'missing': nan,
    'monotone_constraints': None,
    'multi_strategy': None,
    'n_estimators': 100,
    'n_jobs': None,
    'num_parallel_tree': None,
    'random_state': 42,
    'reg_alpha': 0.0,
    'reg_lambda': None,
    'sampling_method': None,
    'scale_pos_weight': None,
    'subsample': 0.5,
    'tree_method': None,
    'validate_parameters': None,
    'verbosity': None,
    'eta': 0.01}

1.3.5 LightGBM

```python
from lightgbm import LGBMRegressor
from sklearn.model_selection import ParameterSampler
```

```python
dict_param = {
    'n_estimator': [10, 20, 50, 100, 200, 500],
    'max_depth': np.asarray([5, 7, 9, 11, 13]).tolist(),
    'num_leaves': ((np.power(2, np.asarray([5, 7, 9, 11, 13])) - 1) * (0.55 +
 ↪(0.65 - 0.55) * np.random.rand(5))).astype(int).tolist(),
    'min_data_in_leaf': np.linspace(100, 1000, 4).astype(int).tolist(),
    'feature_fraction': np.linspace(0.6, 1, 3),
    'bagging_fraction': np.linspace(0.6, 1, 3),
    'learning_rate': [0.01],
    'verbose': [-1],
    'random_state': [42]
}


def custom_sampler(param_grid):
    for params in ParameterSampler(param_grid, n_iter=1e9):
        range_num_leaves = ((0.5 * (2**params['max_depth'] - 1)), (0.7 *
 ↪(2**params['max_depth']) - 1))
        if(range_num_leaves[0] <= params['num_leaves'] <= range_num_leaves[1]):
            for key, value in params.items():
                params[key] = [value]
            yield params
```

```python
grid_search = GridSearchCVTrainer(name='LightGBM Regressor',
 ↪model=LGBMRegressor(),
                                  param_grid=list(custom_sampler(dict_param)), cv
 ↪= 5, n_jobs=2)
grid_search.load_if_exists()
grid_search.fit(X_train, y_train_log)

lgbmr_model = grid_search.best_estimator_
lgbmr_model.fit(X_train, y_train_log)
```

```
c:\Users\aupho\AppData\Local\Programs\Python\Python311\Lib\site-
packages\sklearn\model_selection\_search.py:320: UserWarning: The total space of
parameters 5400 is smaller than n_iter=1000000000. Running 5400 iterations. For
exhaustive searches, use GridSearchCV.
  warnings.warn(

There is no checkpoint file for this model.

100%|        | 1080/1080 [08:50<00:00,  2.03it/s]
```

```python
LGBMRegressor(bagging_fraction=0.6, feature_fraction=0.8, learning_rate=0.01,
              max_depth=11, min_data_in_leaf=100, n_estimator=10,
```

```
                num_leaves=1248, random_state=42, verbose=-1)
```

```
[ ]: evaluate_model(lgbmr_model, 'LightGBM regressor model', X_test, y_test,␣
      ↪y_logscale=True,  save_directory='results/BoW/')
```

```
    LightGBM regressor model's evaluation results:
     - Mean squared error:       3.61
     - Root mean squared error: 1.90
     - Mean absolute error:      1.38
     - R2 error:                 0.02
     - F1 score:                 0.32
     - Precision:                0.26
     - Recall:                   0.41
     - Accuracy:                 0.41
    ------------------------------------------------
```

```
    c:\Users\aupho\AppData\Local\Programs\Python\Python311\Lib\site-
    packages\lightgbm\basic.py:1218: UserWarning: Converting data to scipy sparse
    matrix.
      _log_warning("Converting data to scipy sparse matrix.")
```

```
[ ]: lgbmr_model.get_params()
```

```
[ ]: {'boosting_type': 'gbdt',
      'class_weight': None,
      'colsample_bytree': 1.0,
      'importance_type': 'split',
      'learning_rate': 0.01,
      'max_depth': 11,
      'min_child_samples': 20,
      'min_child_weight': 0.001,
      'min_split_gain': 0.0,
      'n_estimators': 100,
      'n_jobs': None,
      'num_leaves': 1248,
      'objective': None,
      'random_state': 42,
      'reg_alpha': 0.0,
      'reg_lambda': 0.0,
      'subsample': 1.0,
      'subsample_for_bin': 200000,
      'subsample_freq': 0,
      'verbose': -1,
      'n_estimator': 10,
      'min_data_in_leaf': 100,
      'feature_fraction': 0.8,
      'bagging_fraction': 0.6}
```

1.3.6 Stacked model:

```python
from mlxtend.regressor import StackingCVRegressor
```

Define component models:

```python
directory = 'trainer_checkpoint_data/'
with open(directory + 'elastic net_checkpoint.json', 'r') as json_file:
    data = json.load(json_file)
    elastic_model = ElasticNet(**data['best_params'])

with open(directory + 'support vector regressor_checkpoint.json', 'r') as_
 ↪json_file:
    data = json.load(json_file)
    svr_model = SVR(**data['best_params'])

with open(directory + 'random forest regressor_checkpoint.json', 'r') as_
 ↪json_file:
    data = json.load(json_file)
    rfr_model = RandomForestRegressor(**data['best_params'], n_jobs=-1)

with open(directory + 'xgboost regressor_checkpoint.json', 'r') as json_file:
    data = json.load(json_file)
    xgb_model = XGBRegressor(**data['best_params'])

with open(directory + 'lightgbm regressor_checkpoint.json', 'r') as json_file:
    data = json.load(json_file)
    lgbmr_model = LGBMRegressor(**data['best_params'])
```

Define blended model:

```python
stack_gen = StackingCVRegressor(regressors=(xgb_model, lgbmr_model, svr_model,_
 ↪elastic_model, rfr_model),
                                meta_regressor=rfr_model,
                                use_features_in_secondary=True, n_jobs=-1,_
 ↪random_state=42)
stack_gen.fit(X_train, y_train_log)
```

```
StackingCVRegressor(meta_regressor=RandomForestRegressor(max_depth=1000,
                                                         max_features=200,
                                                         min_samples_split=25,
                                                         n_estimators=1024,
                                                         n_jobs=-1,
                                                         random_state=42),
                    n_jobs=-1, random_state=42,
                    regressors=(XGBRegressor(base_score=None, booster=None,
                                             callbacks=None,
                                             colsample_bylevel=None,
```

```
                                      colsample_bynode=None,
                                      colsample_bytree=None, device=None,
                                      early_stopping_rounds…
                                       learning_rate=0.01, max_depth=11,
                                       min_data_in_leaf=100,
                                       n_estimator=10, num_leaves=1248,
                                       random_state=42, verbose=-1),
                       SVR(C=1000, gamma=1e-06),
                       ElasticNet(alpha=100, l1_ratio=0.4,
                                  max_iter=2000, random_state=42),
                       RandomForestRegressor(max_depth=1000,
                                             max_features=200,
                                             min_samples_split=25,
                                             n_estimators=1024,
                                             n_jobs=-1,
                                             random_state=42)),
           use_features_in_secondary=True)
```

```
[ ]: evaluate_model(stack_gen, 'Stacking model', X_test, y_test, y_logscale=True,␣
     ↪save_directory='results/BoW/')
```

c:\Users\aupho\AppData\Local\Programs\Python\Python311\Lib\site-
packages\lightgbm\basic.py:1218: UserWarning: Converting data to scipy sparse
matrix.
  _log_warning("Converting data to scipy sparse matrix.")

```
Stacking model's evaluation results:
 - Mean squared error:      3.36
 - Root mean squared error: 1.83
 - Mean absolute error:     1.38
 - R2 error:                0.09
 - F1 score:                0.29
 - Precision:               0.34
 - Recall:                  0.28
 - Accuracy:                0.28
--------------------------------------------------
```