

# Storypoint Problem Exploration - usergrid

September 6, 2024

## 1 Storypoint Prediction: Problem Exploration

### 1.1 Problem Statement

In modern agile development settings, software is developed through repeated cycles (iterative) and in smaller parts at a time (incremental), allowing for adaptation to changing requirements at any point during a project's life. A project has a number of iterations (e.g. sprints in Scrum). Each iteration requires the completion of a number of user stories, which are a common way for agile teams to express user requirements.

There is thus a need to focus on estimating the effort of completing a single user story at a time rather than the entire project. In fact, it has now become a common practice for agile teams to go through each user story and estimate its "size". Story points are commonly used as a unit of measure for specifying the overall size of a user story.

### 1.2 Problem Formulation

**Input:** A string of length  $N$  that contains a story's name and description  $C = \{c_1, c_2, c_3, \dots, c_n\}$ . For each story, a set of text embeddings that contains features  $E = \{e_1, e_2, e_3, \dots, e_m\}$  extracted from  $C$  has been provided.

**Output:** A natural number  $P$  associated with the story point of that user story

### 1.3 Dataset Information

**Text Embeddings:** Text embeddings are a way to convert words or phrases from text into a list of numbers, where each number captures a part of the text's meaning. The dataset has been preprocessed and converted into two kinds of text embeddings. You can choose to work with either of them or both: - **Doc2Vec:** Input strings are transformed into fixed-length vectors of size 128. These vectors capture the semantic meaning of words and their relationships within a document. - **Look-upTable:** Input strings are transformed into fixed-length vectors of size 2264. These vectors are obtained via transforming each word in the input strings into an identifier number, then padded to the length of the longest sample.

**Dataset Structure & Format:** Storypoint Estimation Dataset is stored in 3 folders labeled *raw data*, *look-up*, and *doc2vec*. Within each folder are 3 CSV files for training, testing, validation. Each csv file has the following columns: - **issuekey** : The unique identifier for a story. - **storypoint**: The correct number of storypoint. - An embedding column (**embedding** or **doc2vec**) contains text embedding vectors. The raw data csv will not have this and instead contain two columns with **story name** and **description**.

## 1.4 Exploration

### 1.4.1 Raw data exploration

```
[ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.feature_extraction.text import CountVectorizer
```

#### Output exploration

```
[ ]: # Import raw data from the CSV file

project_name = 'usergrid'

all_data = pd.concat([pd.read_csv('data/' + project_name + '/' + project_name +
    ↪ '_train.csv'),
                    pd.read_csv('data/' + project_name + '/' + project_name +
    ↪ '_valid.csv'),
                    pd.read_csv('data/' + project_name + '/' + project_name +
    ↪ '_test.csv')])

print('Check the shape of the dataset', all_data.shape)
```

Check the shape of the dataset (333, 4)

```
[ ]: all_data.drop(['issuekey'], axis=1, inplace=True)
all_data.head()
```

```
[ ]:                                     title \
0  asset data correctly obey contextual ownership...
1                                     expose refresh token rest tier
2                      bad geo query returns entire collection
3  empty string doesnt remove entity property nul...
4  fresh admin user token wont work managementuse...

                                     description  storypoint
0  asset data endpoint assetsuiddata correctly o...          3
1  need add refresh token capability rest tier ma...          3
2  badly formed geo query sent get api returns en...          3
3  users able remove entity properties using eith...          3
4  logging curl post granttypepasswordusernamefds...          3
```

First, let take a look at the distribution of the story point:

Interpretation of Skewness Values:

- **Skewness > 0:** Right-skewed distribution.
- **Skewness < 0:** Left-skewed distribution.

- **Skewness = 0**: Symmetrical distribution (like a normal distribution).


Interpretation of kurtosis: - **Leptokurtic (Kurtosis > 3)**: The distribution has heavier tails and a sharper peak than the normal distribution. Data points are more likely to produce extreme values. The distribution has a higher peak and fatter tails. - **Platykurtic (Kurtosis < 3)**: The distribution has lighter tails and a flatter peak than the normal distribution. Data are fewer extreme values compared to a normal distribution. - **Mesokurtic (Kurtosis = 3)**: The distribution has a similar kurtosis to the normal distribution, indicating a moderate level of outliers.

```
[ ]: # Draw a histogram of the story points
plt.figure(figsize=(12, 4))
plt.xticks(np.arange(0, max(all_data['storypoint']) + 1, 5))
sns.histplot(all_data['storypoint'], bins=100, kde=True)

print('Skewness:', all_data['storypoint'].skew())
print('Kurtosis:', all_data['storypoint'].kurt())
```

Skewness: 1.2353032737625909

Kurtosis: 2.7202275525441997



```
[ ]: tmp = pd.concat([all_data['storypoint'].value_counts(),
                      all_data['storypoint'].value_counts() / all_data.shape[0] * 100],
                      axis=1, keys=['Counts', 'Percentage (%)'])
tmp.head(20)
```

```
[ ]:
storypoint
3          189      56.756757
5           48      14.414414
1           47      14.114114
2           38      11.411411
8           11       3.303303
```


This data is a little bit right-skewed and separate in 5 groups. Maybe this problem is suitable for a classification method than a regression method

I will try 2 solutions: - Use log-scale on the label - Remove all the examples with label greater than a threshold (20, 30 or 40)

The first solution: logarithm magic

```
[ ]: sns.histplot(np.log(all_data['storypoint']), bins=20, kde=True)
```

```
[ ]: <Axes: xlabel='storypoint', ylabel='Count'>
```



```
[ ]: print('Skewness:', np.log(all_data['storypoint']).skew())  
print('Kurtosis:', np.log(all_data['storypoint']).kurt())
```

Skewness: -0.5822367330466891

Kurtosis: 0.3649582125623412

Kurtosis now becomes platokurtosis but skewness is better than before

The second solution: Dismantle and Cleave

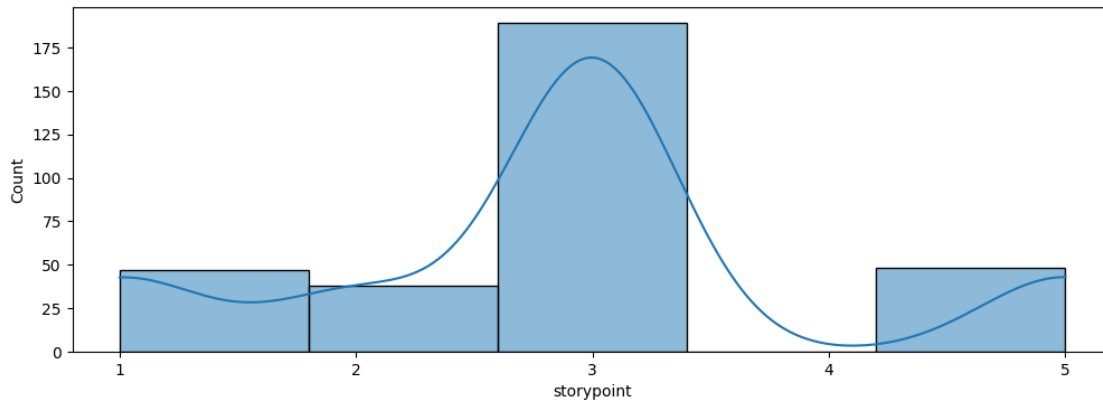
```
[ ]: threshold = 5 # This threshold means that we will take all the examples with  
    ↪ story points less than or equal to 5
```

```

new_data = all_data[all_data['storypoint'] <= threshold]
plt.figure(figsize=(12, 4))
plt.xticks(np.arange(0, max(new_data['storypoint']) + 1, 1))
sns.histplot(new_data['storypoint'], bins=threshold, kde=True)
print('Fitered percentage: ', round(1 - new_data.shape[0] / all_data.shape[0],
↪2) * 100, '%')

```

Fitered percentage: 3.0 %



**Input exploration** The input of this problem is 2 texts: title and description. First we will find some statistics:

```

[ ]: title_lengths = all_data['title'].apply(lambda x: len(x.split(' ')))
print('Title analysis:')
print('  - Mean length:', round(title_lengths.mean()))
print('  - Min length:', title_lengths.min())
print('  - Max length:', title_lengths.max())

description_lengths = all_data['description'].apply(lambda x: len(x.split(' '))
↪if type(x) != float else 0)
print('Description analysis:')
print('  - Mean length:', round(description_lengths.mean()))
print('  - Min length:', description_lengths.min())
print('  - Max length:', description_lengths.max())

```

Title analysis:

- Mean length: 6
- Min length: 1
- Max length: 13

Description analysis:

- Mean length: 36
- Min length: 0

- Max length: 436


Plot the histogram of the title length and KDE of the description length (exclude 0):

```
[ ]: plt.figure(figsize=(12, 8))

plt.subplot(2, 1, 1)
plt.xticks(np.arange(0, max(title_lengths) + 1, 1))
plt.xlabel('Title Length')
sns.histplot(title_lengths, bins=max(title_lengths))

plt.subplot(2, 1, 2)
plt.xlabel('Description Length')
plt.xscale('log')
sns.kdeplot(description_lengths[description_lengths > 0])
```

```
[ ]: <Axes: xlabel='Description Length', ylabel='Density'>
```



I think we should check the correlation between title length and description length:


```
[ ]: plt.figure(figsize=(12, 4))

plt.subplot(1, 2, 1)
```

```
plt.xticks(np.arange(0, max(all_data['title'].apply(lambda x: len(x.split(' ')))) + 1, 5))
sns.regplot(x=all_data['title'].apply(lambda x : len(x.split(' '))),
            y=all_data['storypoint'],
            line_kws={'color': 'red'})

plt.subplot(1, 2, 2)
sns.regplot(x=all_data['description'].apply(lambda x : len(x.split(' ')) if
↪type(x) != float else 0),
            y=all_data['storypoint'],
            line_kws={'color': 'red'})
```

```
[ ]: <Axes: xlabel='description', ylabel='storypoint'>
```



The correlation is having too much deviation on the right plot and a little bit on the left.

Let dive deeper in the input:

Title analysis:

```
[ ]: count_vectorizer = CountVectorizer()
count_vectorizer.fit(all_data['title'])

dictionary = pd.DataFrame(list(count_vectorizer.vocabulary_.items()),
↪columns=['word', 'frequency'])
dictionary.sort_values(by='frequency', ascending=False, inplace=True)
print(dictionary.shape)
dictionary.head(10)
```

```
(828, 2)
```

```
[ ]:      word  frequency
138  wrong      827
232  writes      826
```

381	writereview	825
403	write	824
227	works	823
45	working	822
170	workflow	821
25	work	820
30	wont	819
665	without	818

Description analysis:

```
[ ]: count_vectorizer = CountVectorizer()
count_vectorizer.fit(all_data[all_data['description'].isnull() ==
↪False]['description'])

dictionary = pd.DataFrame(list(count_vectorizer.vocabulary_.items()),
↪columns=['word', 'frequency'])
dictionary.sort_values(by='frequency', ascending=False, inplace=True)
print(dictionary.shape)
dictionary.head(20)
```

(3003, 2)

```
[ ]:
      word frequency
749      zookeeper      3002
2807      zip      3001
119  ywmtxkhayeeoblzqpxkkaaaudriyfhegizloihjztoxj...      3000
1151 ywmtslyneesmfpujajlwaaauvjenpfneqwydblynuxagsh...      2999
604  ywmtmzwuiqueescltpguhfegeaaufjuwrktiypwsnvohbv...      2998
608  ywmtjjacqleesrfvxwfraaaauffucqvwoyiavxxioeafnoa      2997
113  ywmtifzejreeohfgowhrloaaaauizmdwtwqtdewhckgdy...      2996
2560      yields      2995
498      yet      2994
2745      yes      2993
2282      yattributeaction      2992
1717      xyz      2991
539      xxx      2990
517      xsd      2989
2281      xobjectconcept      2988
1259      xfddd      2987
1286      xfcbc      2986
1263      xeeded      2985
1287      xedccf      2984
1280      xeda      2983
```

Yet I don't find any thing special about the words in input except so many things are bad.



### 1.4.2 Solving strategies

My first intuition in this problem is that the hard part is not on the algorithm we use, it is on the **embedding** part. Therefore, in case the given embedded datasets work not properly, I will use a better embedding method which is **Bidirectional Encoder Representations from Transformers (BERT)**. Also, I will try an old way to embedding the text too: **Bag of words**.

In conclusion, I will have 4 ways to embed the text: - doc2vec (already available) - Look up (already available) - Bag Of Words - BERT

About algorithm, I will try all the regression algorithm that may give a good result:

- Ridge Regressor
- Support Vector Regressor
- Random Forest Regressor
- Gradient Boosting
- XGBoost
- Lightgbm
- Blended

Maybe, we can change the problem to the classification problem with 100 labels (desparation confirmed). In the classification problem, I will use: - Support Vector Classifier - Softmax Regression (Multinomial Logistic Regression) - Random Forest - Adaboost - XGBoost

Thanks to the libraries, the implementation of all the algorithm shrinks to its minimum form.

At last, there is still a situation that all of mentioned model don't give a good result. This gamble is thrilling (hopeless).

*"But would you lose?"*

Nah, I'd win.